Housing price modelling competition

Leonardo Blas

# Can we predict housing prices?

- We have:
  - A dataset of properties sold from 2006 to 2010 in Ames, IA.
  - ~1500 entries.
  - ~80 parameters.
- We want:
  - A model to predict housing prices.
  - To minimize the error in terms of dollars (RMSE).

Process overview

1. Consider the features we want to use.

2. Clean the data.

3. Transform the data (encode, standardize, normalize).

4. Try different models.

5. Choose the best performing model.

What features to use

- Overall quality.
- Living area area.
- Basement area.
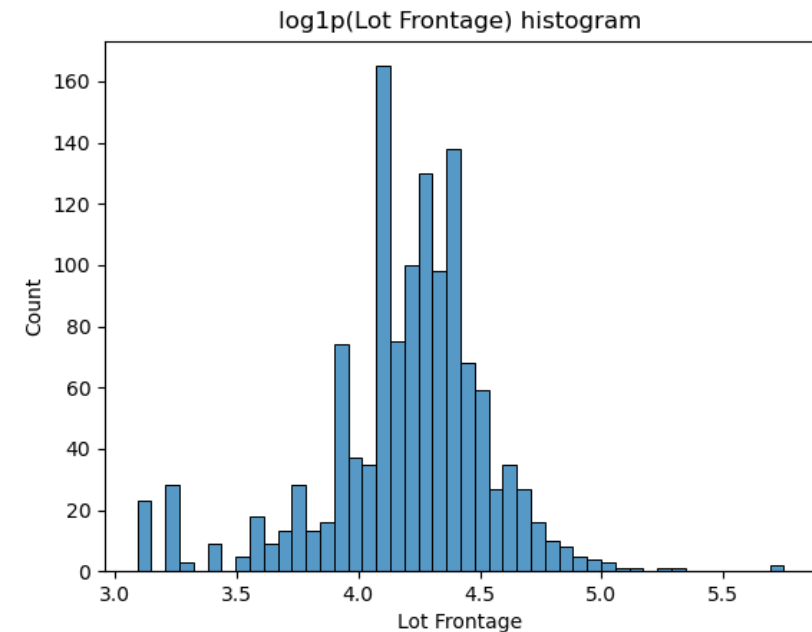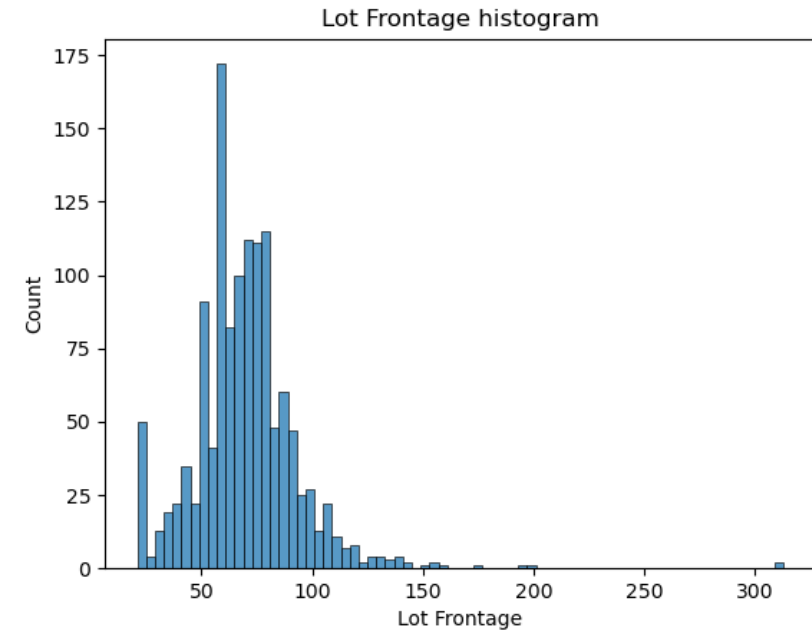- Garage car capacity.
- Garage area.

Features with low correlation to SalePrice

What not to use

- Overall condition.
- Year sold.
- Month sold.
- Pool area.
- Basement half bathrooms.

Features with low correlation to SalePrice

Helpful transformations

- log1p.
- Power.
- KNN imputation.
- Standard scaling.

Lot Frontage histogram

log1p(Lot Frontage) histogram

My models: Linear

- Train RMSE: ~$17k
- Test RMSE: ~$21k

- Train r^2: ~0.95
- Test r^2: ~0.93

Linear model: Predicted price vs actual price

My models: Ridge

- Train MSE: ~$21k
- Test MSE: ~22k
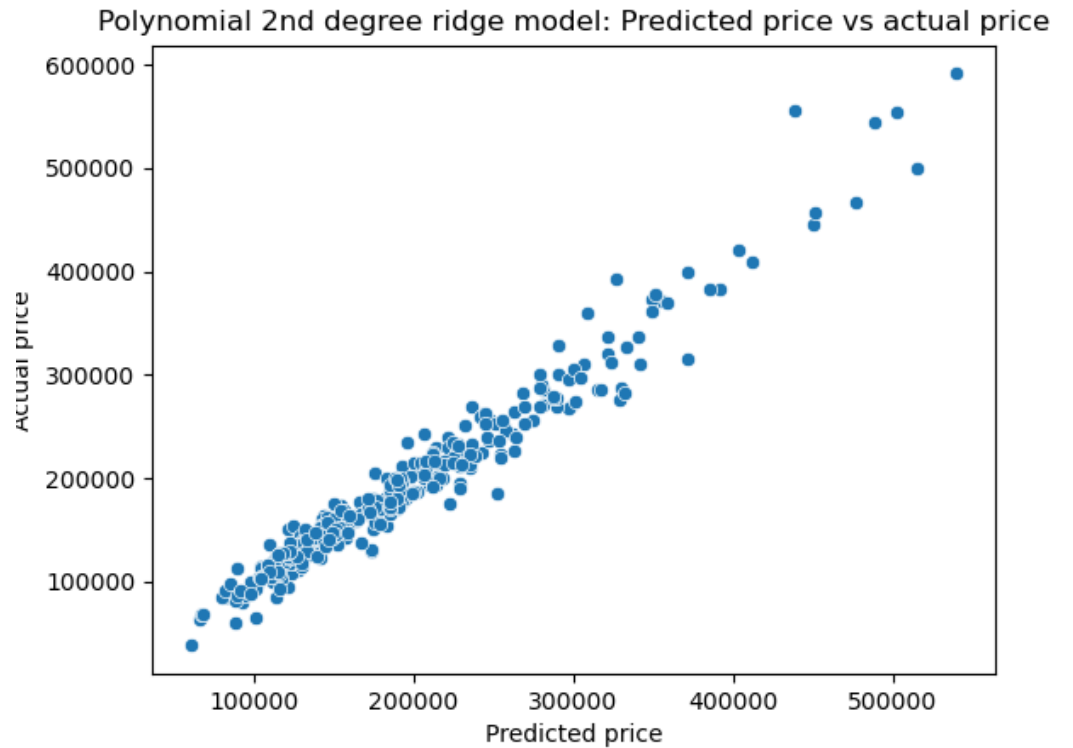
- Train r^2: 0.94
- Test r^2: 0.93

Ridge model: Predicted price vs actual price

My models: Polynomial linear

- Train RMSE: ~$34k
- Test RMSE: ~$68k

- Train r^2: 1.0
- Test r^2: 0.87

Polynomial 2nd degree linear model: Predicted price vs actual price

My models: Polynomial ridge

- Train RMSE: ~$19k
- Test RMSE: ~$17k

- Train r^2: ~0.97
- Test r^2: ~0.96

Polynomial 2nd degree ridge model: Predicted price vs actual price

Findings and recommendations

- Data doesn't lie: Build models considering the data patterns.

- It's hard to summarize inter-parameter interactions.

- Parameters like overall quality and garage area are important.

- Some data is almost irrelevant.

- A 2^{nd} degree polynomial ridge model is efficient.

References

- [1] J. O, M. Harris, "Ames Iowa Submission". 2024. https://kaggle.com/competitions/adobe-dsb-34.